




# **Lecture 2: MPSoC intro, ZYNQ chip, basic design, and debugging**

(UART with PWM)

Emad Samuel Malki Ebeid,  
University of Southern Denmark

# Is it **FPGA**, **CPU**, or **GPU** to us in autonomous robots?



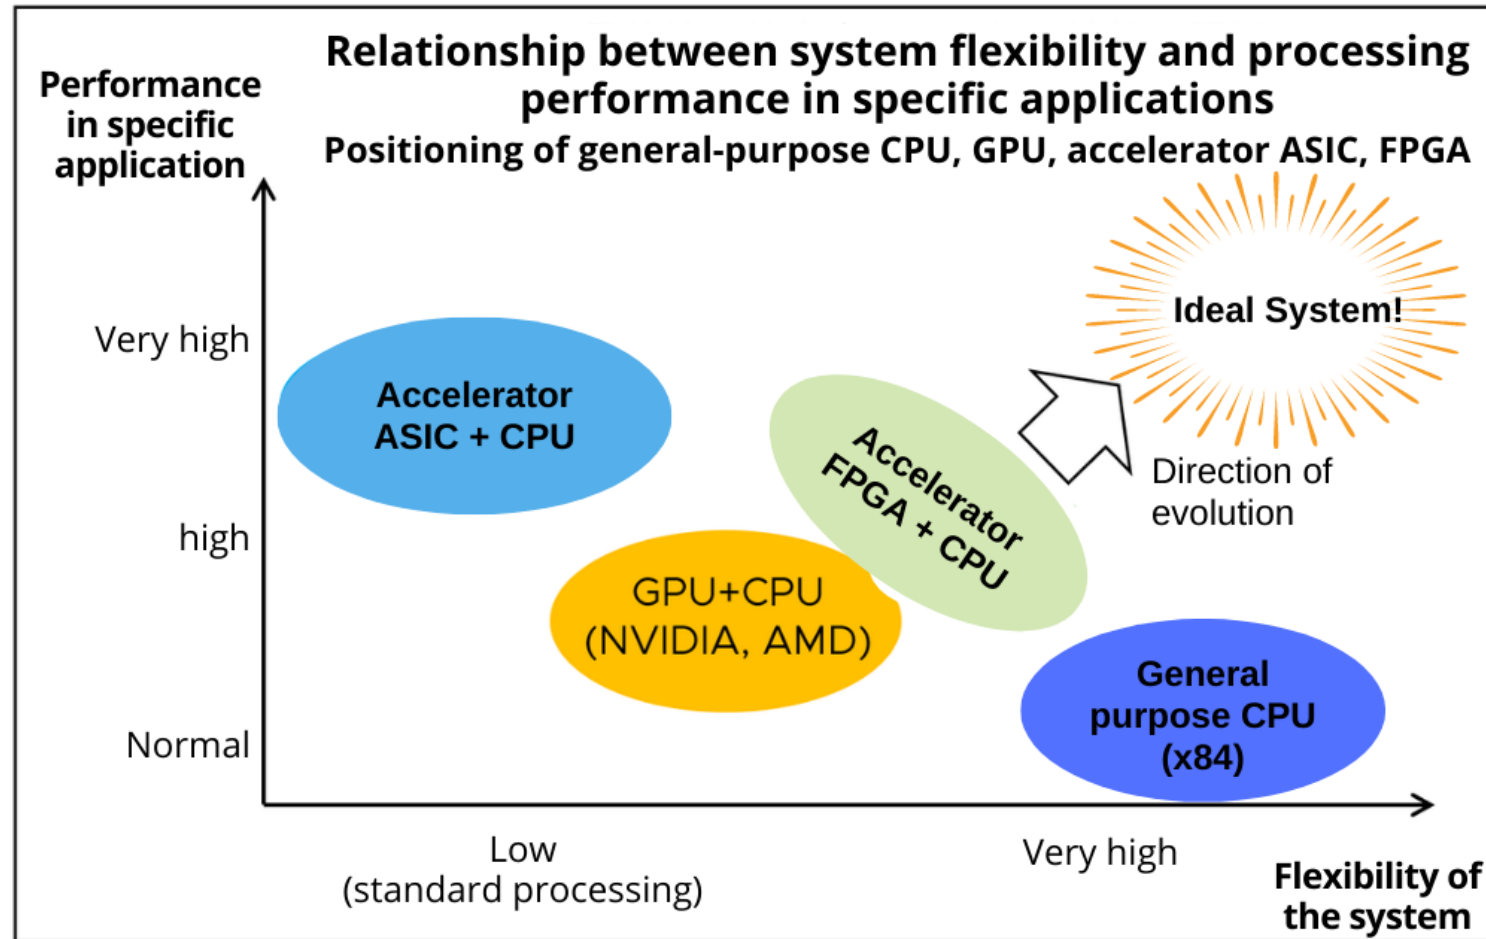
# Introduction

- Autonomous operations of drones have significantly increased in the last few years, spanning from autonomous infrastructure inspections to interactions and manipulation of hard-to-reach objects.
- Robots and Drones integrate several technologies and algorithms for accomplishing their tasks (e.g., path planning and scene reconstruction, 3D sensing, localization, exploration and navigation). That poses many challenges due to the enormous amount of data that needs to be processed onboard of the drone in real time. Commonly, CPUs and GPUs are used to develop novel algorithms and handle a wide range of tasks quickly.

# Real-time system requirements?

- Algorithms are typically captured in C/C++ or some other high-level language, which abstracts the details of the computing platform.
- Improving the runtime of software was based on two central concepts:
  - increasing processor clock frequency and
  - using specialized processors
- Incremental speedup through processor clock frequency is not enough to deliver a viable product to market.
- Speedup by adding more processing cores per chip.
- Multicore processors to boost software performance.

# Performance



[https://www.xilinx.com/support/documentation/sw\\_manuals/ug998-vivado-intro-fpga-design-hls.pdf](https://www.xilinx.com/support/documentation/sw_manuals/ug998-vivado-intro-fpga-design-hls.pdf)

<https://hacarus.com/tech/20201207-amd-acquires-xilinx/>

# Performance

- **Clock Frequency:** a high clock frequency translates into a higher performance execution rate of an algorithm?
- Instruction fetch (IF)
- Instruction decode (ID)
- Execute (EXE)
- Memory operations (MEM)
- Write back (WB)




Figure 3-1: Processor Instruction Execution Stages

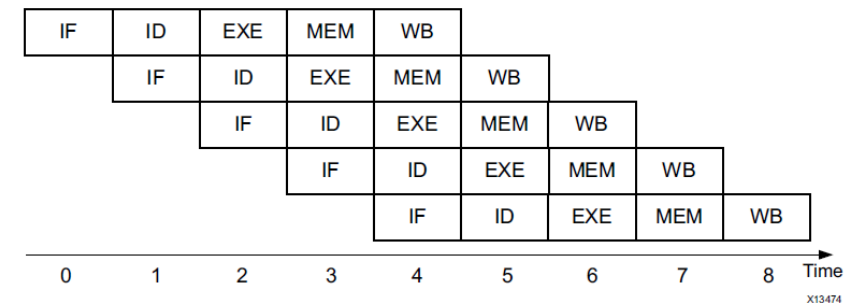


Figure 3-2: Processor with Multiple Instruction Execution Units

# Performance

- FPGA executes a single program at a time on a custom circuit for that program.
- Therefore, changing the user application changes the circuit in the FPGA.




Figure 3-3: FPGA Instruction Execution Stages




Figure 3-4: FPGA with Multiple Instruction Execution Units

- the FPGA has a nominal performance advantage of 9x compared to the processor

# Performance metrics

- **Power consumption**

- $P = 0.5 * C * V^2 * f$

- Where P is power, C is capacitance, V is the voltage across the gate, f is the clock frequency.

- By creating a custom circuit per software program, an FPGA is able to run at a lower clock frequency with maximum parallelism between operations and without the instruction interpretation overhead found in a processor.



# Latency

- **Latency** is the number of clock cycles it takes to complete an instruction or set of instructions to generate an application result value
  - five clock cycles/instruction.




Figure 3-1: Processor Instruction Execution Stages

- Pipelining to lower the latency: the next instruction can be launched into execution before the current instruction is complete

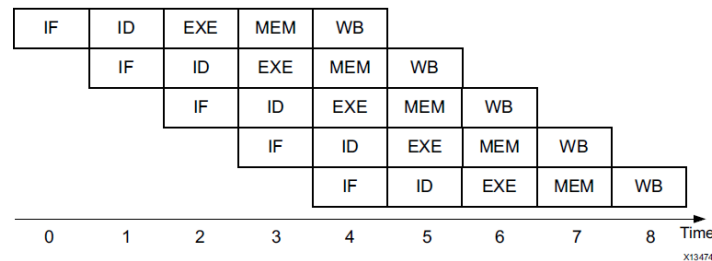


Figure 3-2: Processor with Multiple Instruction Execution Units

# Throughput

- **Throughput** is the number of clock cycles it takes for the processing logic to accept the next input data sample.
- Figure 3-5 requires 10 ns between input samples,
- Figure 3-6 only requires 2 ns between input data samples.
- Figure 3-6 circuit implementation has higher performance, because it can accept a higher input data rate.

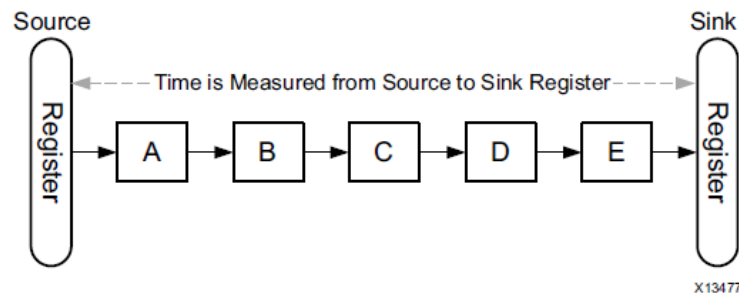


Figure 3-5: FPGA Implementation without Pipelining

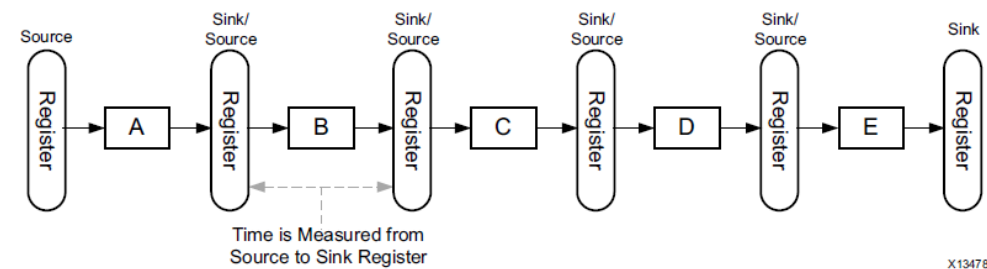


Figure 3-6: FPGA Implementation with Pipelining

# FPGA


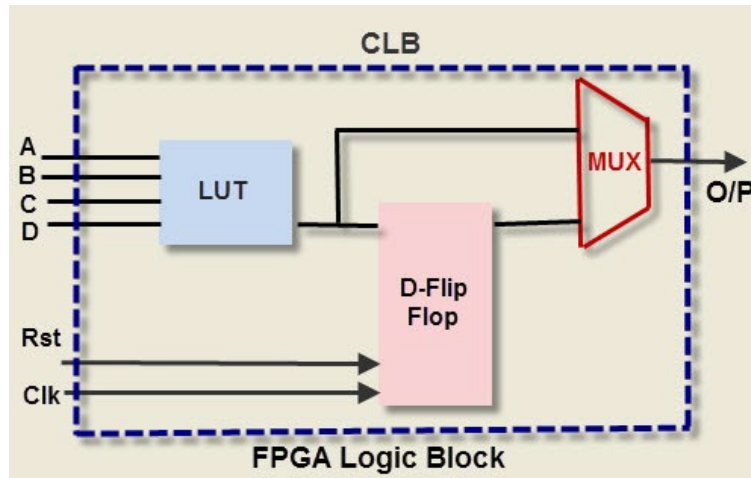
Background knowledge

# FPGA

- Field Programmable Gate Array (FPGA) is an integrated circuit containing gate matrix which can be programmed by the user “in the field” without using expensive equipment.
- An FPGA contains a set of programmable logic gates and rich interconnect resources, making it possible to implement complex digital circuits.
- FPGA devices are produced by a number of semiconductor companies:
  - Xilinx, Altera, Actel, Lattice, QuickLogic and Atmel.
- Configuration bitstream can be stored in FPGA using various technologies.
  - The majority of FPGAs is based on SRAM (Static RAM).

# FPGA

- The general FPGA architecture consists of three types of modules. They are I/O blocks or Pads, Switch Matrix/ Interconnection Wires and Configurable Logic Blocks (CLB).




# Memory cells: RAM

RAM is a **volatile** memory that the data is eventually lost when the memory is not powered




## SRAM: Static random-access memory

- Speed ↑ size (bigger (6 MOSFETs))
- Cache memory
- Low power consumption
- 3 Operation States:
  - hold ( $WL=0$ ), write & read ( $WL=1$ )



## DRAM: Dynamic random-access memory

- Speed ↓ size (smaller (one MOSFET))
- Main memory
- Charge leakage
- High power consumption
- Refreshed frequently (few milliseconds)





Single Memory Cell

# SRAM-based FPGAs

- It stores logic cells configuration data in the static memory
- Since SRAM is volatile, the FPGAs must be programmed (configured) upon start
- There are two basic modes of programming:
  - **Master mode**, when FPGA reads configuration data from an external source, such as an external Flash memory chip.
  - **Slave mode**, when FPGA is configured by an external master device, such as a processor (via a dedicated configuration interface or via a boundary-scan (JTAG) interface).
- SRAM-based FPGAs with an internal flash memory:
  - it contains internal flash memory blocks, thus eliminating the need to have an external non-volatile memory. It uses flash only during startup to load data to the SRAM configuration cells.

# Programmable switch FPGAs

- It uses flash as a primary resource for configuration storage, and doesn't require SRAM
- Switch is a floating-gate transistor that can be turned off by injecting charge onto its floating gate
- It has a limit number of reprogramming
- More secure than SRAM-FPGA







# Fuse-based FPGA


- One-time program (OTP) FPGA
- The fuse makes or breaks the link between two wires
  - The amorphous silicon (nonconductive material) turns into a polycrystalline silicon-metal alloy (conductive) when exposed to high voltage.
- More secure than SRAM (no load from an external device )
- In high radiation environments (space or nuclear reactors), radiation events can cause SRAM, which contains the program, to change state but not in fused FPGA.

# Programmable switch matrix

- All internal connections are composed of metal segments with programmable switching points and switching matrices to implement the desired routing.



# Example: FPGA programming



$$f_1 = x_1x_2$$

$$f_2 = \overline{x_2}x_3$$

$$f = x_1x_2 + \overline{x_2}x_3$$

## Other FPGA Building Blocks

- Clock distribution: element clock skew
- Embedded memory blocks
- Special purpose blocks:
  - DSP blocks: Hardware multipliers, adders and registers
  - Embedded microprocessors/microcontrollers
  - High-speed serial transceivers

# Design flow

Design and implement a simple unit permitting to speed up encryption with RC5-similar cipher with fixed key set on 80331 microcontroller. Unlike in the experiment 5, this time your unit has to be able to perform an encryption algorithm by itself, executing 32 rounds.....



Specification / Pseudocode

On-paper hardware design




```
Library IEEE;
use ieee.std_logic_1164.all;
use ieee.std_logic_unsigned.all;

entity RC5_core is
    port(
        clock, reset, over_done: in std_logic;
        data_input: in std_logic_vector(7 downto 0);
        data_output: out std_logic_vector(7 downto 0);
        out_full: in std_logic;
        key_input: in std_logic_vector(31 downto 0);
        key_read: out std_logic;
        k: out std_logic_vector(31 downto 0);
    );
end RC5_core;
```




Synthesis




VHDL description (Your Source Files)


Functional simulation




Post-synthesis simulation



Implementation




Timing simulation



Configuration



On chip testing

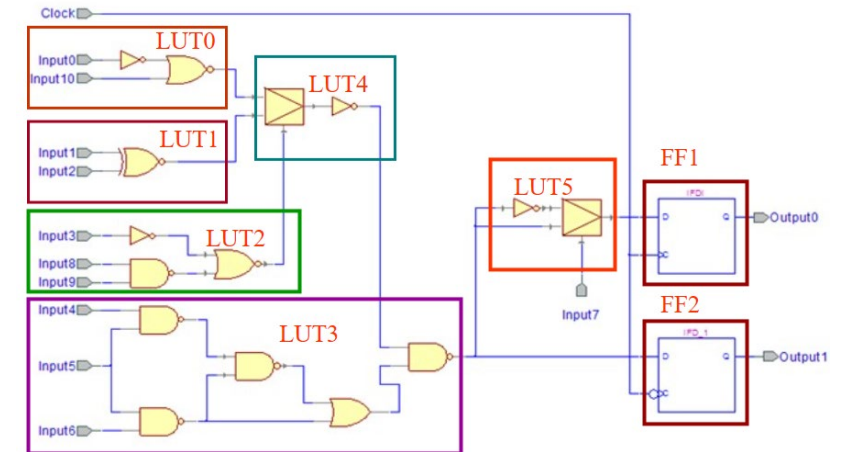



# Synthesis

```
architecture MLU_DATAFLOW of MLU is
    signal A1:STD_LOGIC;
    signal B1:STD_LOGIC;
    signal Y1:STD_LOGIC;
    signal MUX_0, MUX_1, MUX_2, MUX_3: STD_LOGIC;
begin
    A1<=A when (NEG_A='0') else
        not A;
    B1<=B when (NEG_B='0') else
        not B;
    Y<=Y1 when (NEG_Y='0') else
        not Y1;

    MUX_0<=A1 and B1;
    MUX_1<=A1 or B1;
    MUX_2<=A1 xor B1;
    MUX_3<=A1 xnor B1;

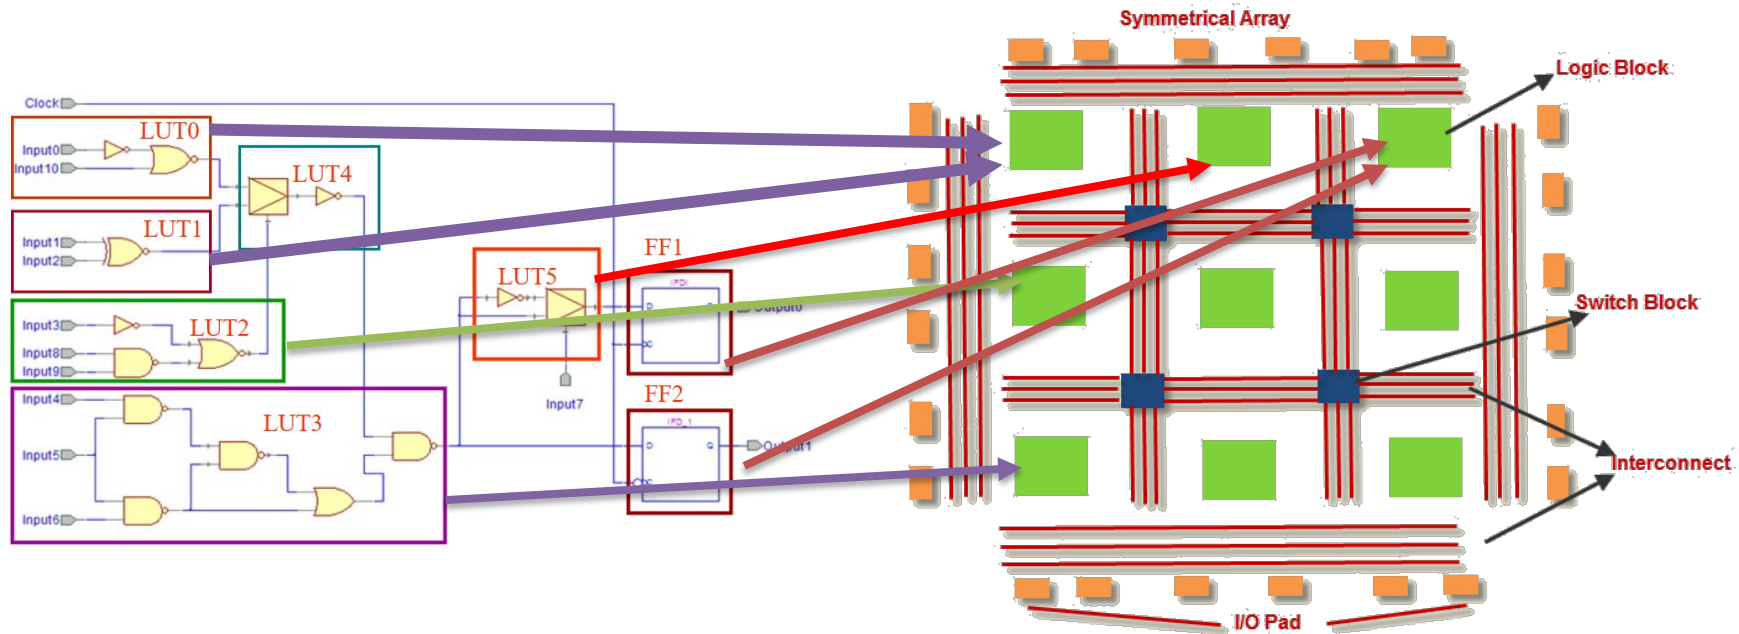
    with (L1 & L0) select
        Y1<=MUX_0 when "00",
            MUX_1 when "01",
            MUX_2 when "10",
            MUX_3 when others;
end MLU_DATAFLOW;
```



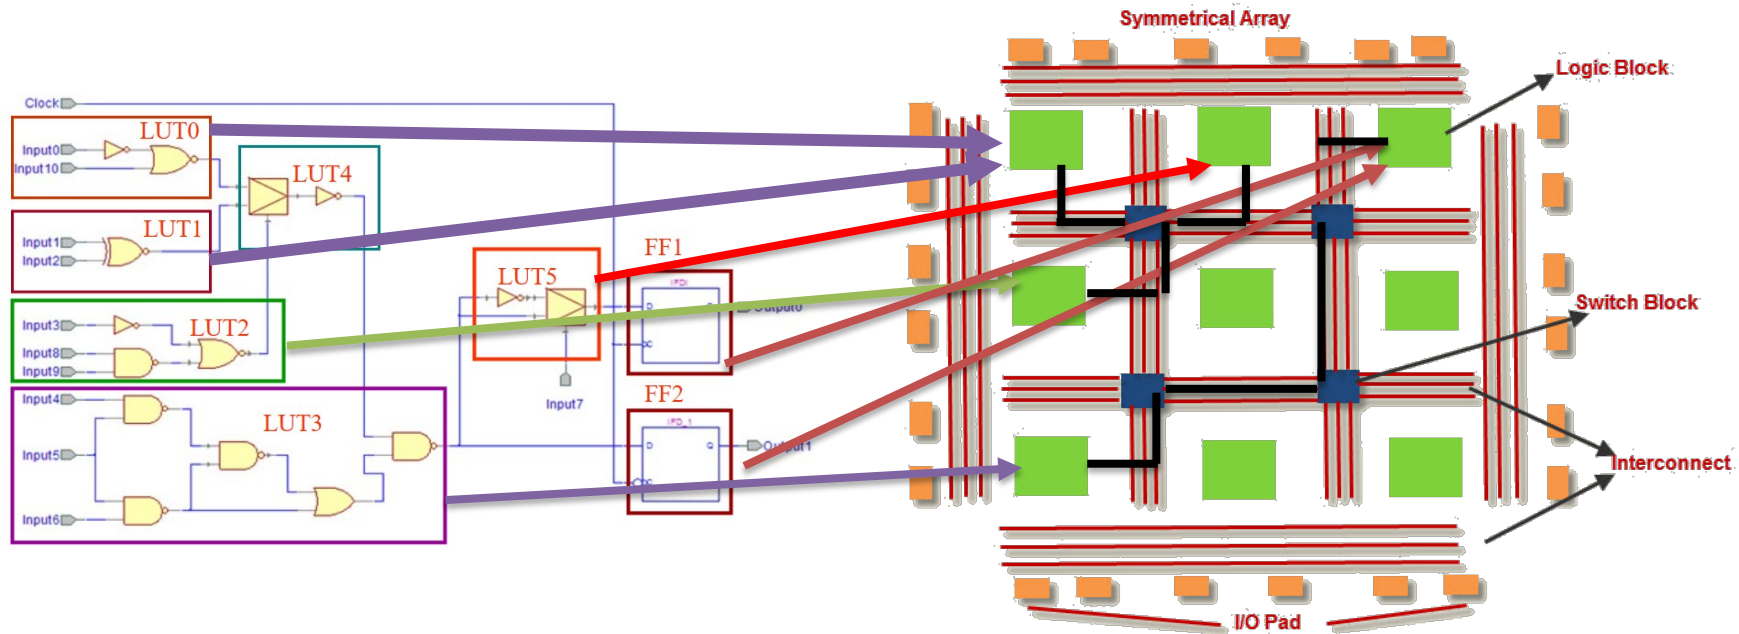
Mapping

Circuit netlist

# Placing



# Routing





# Configuration

- Once a design is implemented, the synthesis tool generates a file that the FPGA can understand
- This file is called a bit stream: a BIT file (.bit extension)
- The BIT file can be downloaded directly to the FPGA, or can be converted into a PROM file which stores the programming information

# **ZYNQ architecture**

## **CPU+FPGA**

# ZYNQ Architecture

- **Zynq** architecture created by Xilinx.
- It combines an FPGA with ARM cores and I/O into one product.
- The ARM part is called **Processing System (PS)**.
  - application processor unit (APU)
  - fully integrated memory controllers
  - I/O peripherals
- The FPGA is called **Programmable Logic (PL)**.

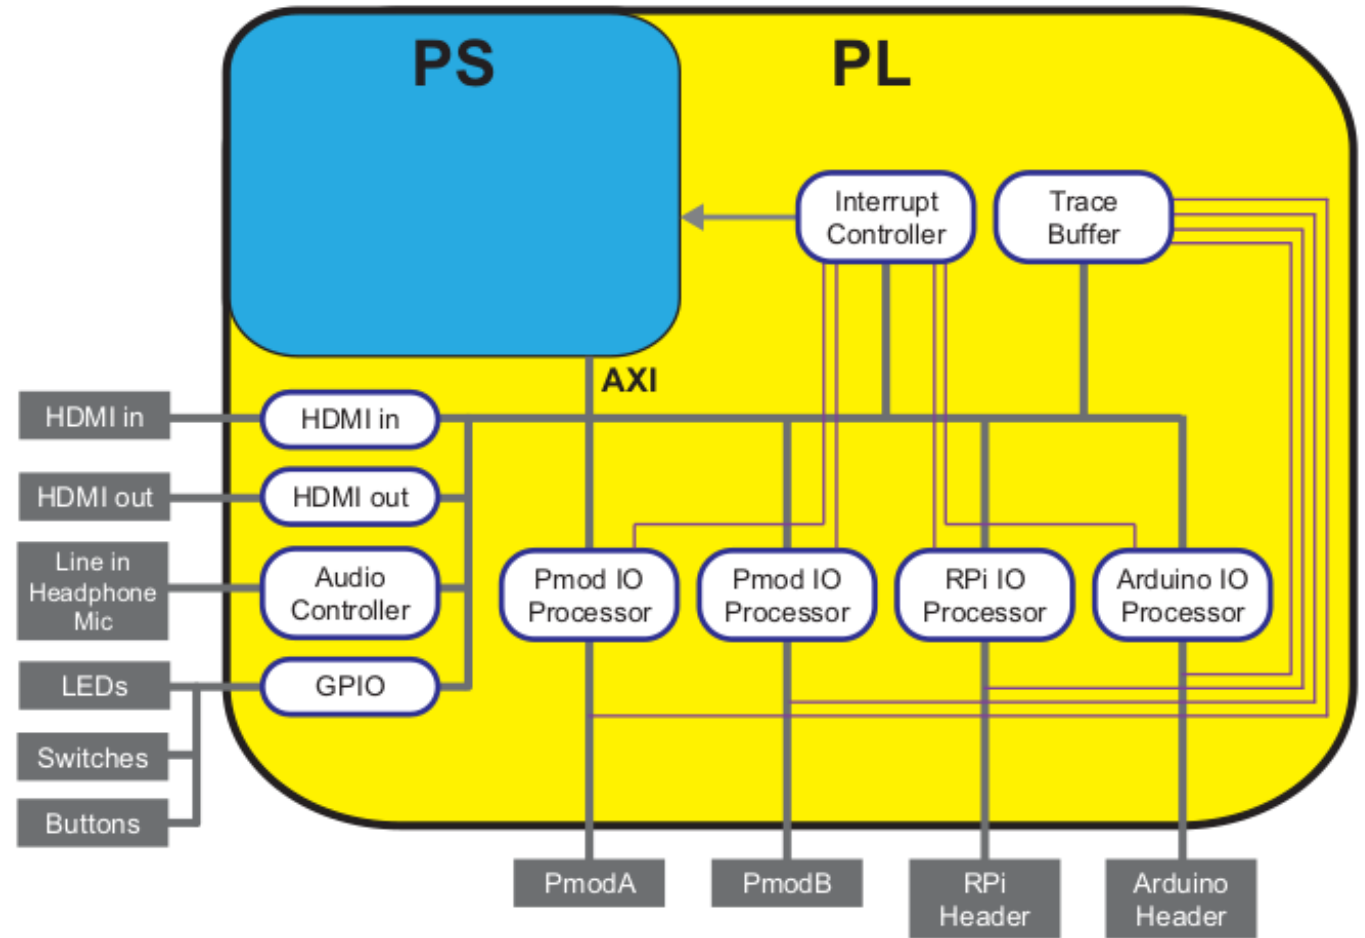


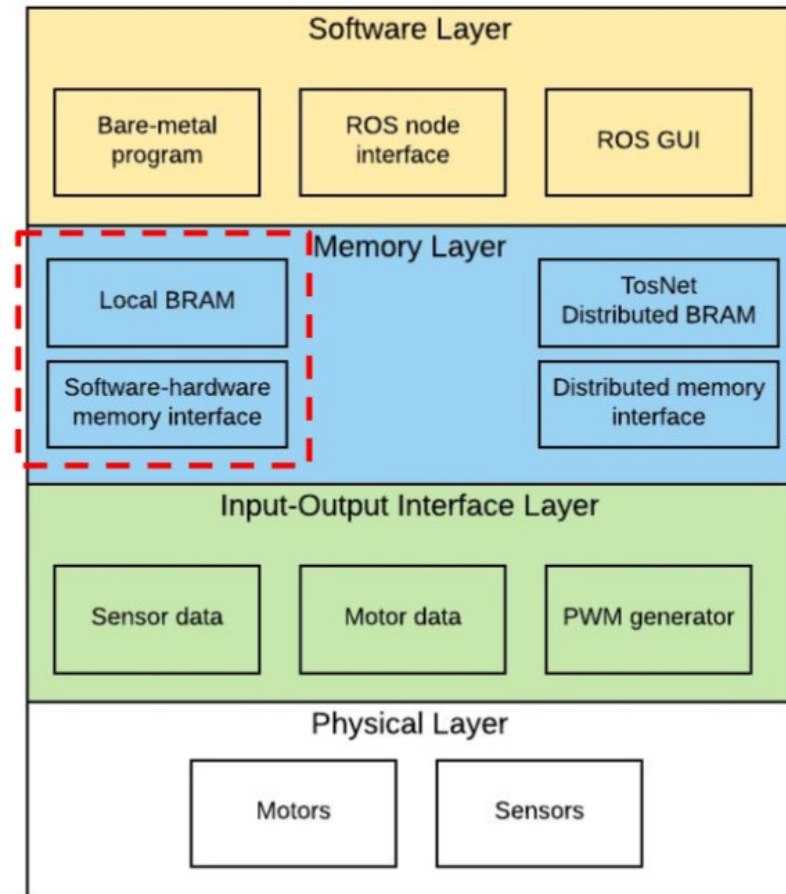


Figure 22.7: Interface functionality implemented by the base overlay (for PYNQ-Z2)

# PS-PL specs


- PS: processing system
  - dual ARM Cortex-A9 processors, 866MHz to 1 GHz frequency
  - multiple peripherals
  - hard silicon core
- PL: programmable logic
  - shares the same FPGA series 7 programmable logic
  - logic cells: 28k - 444k (430k to 6.6M gates)
  - flip-flops: 35k - 554k
  - DSP/MAC: 80 - 2020
  - AD converter: two 12bits

# Example Implementation



[http://weisong.eng.wayne.edu/\\_resources/pdfs/wu20-HydraMini.pdf](http://weisong.eng.wayne.edu/_resources/pdfs/wu20-HydraMini.pdf)

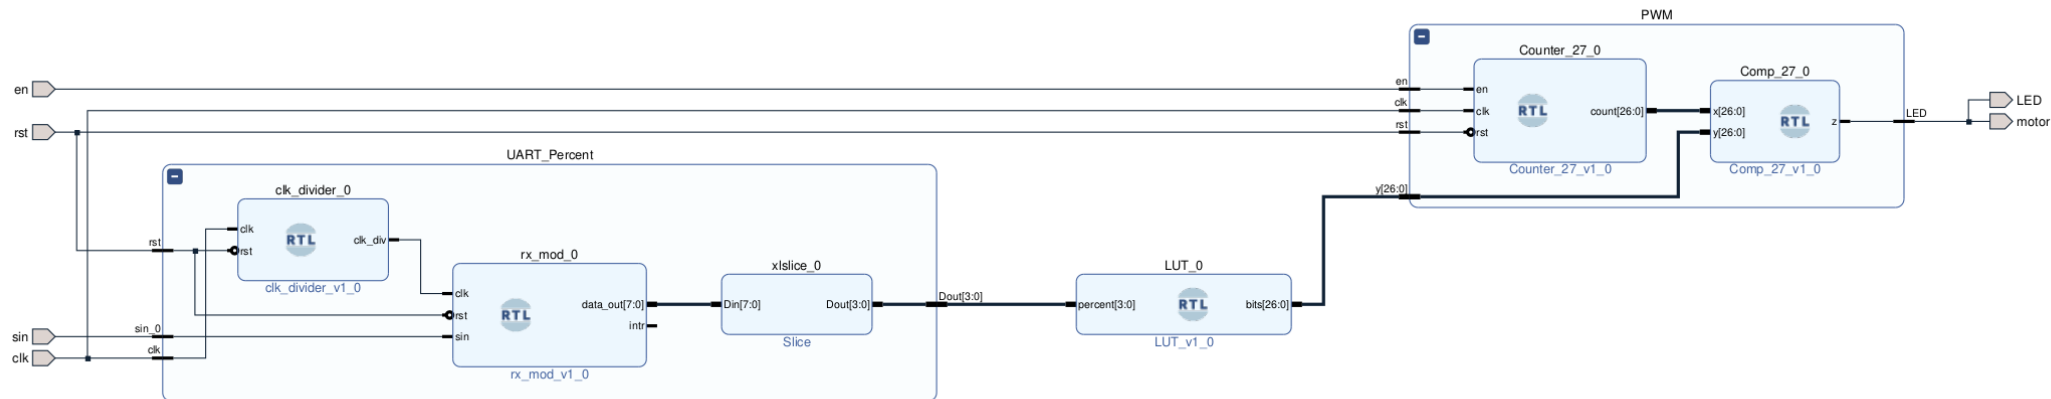
# MPSoC



# Lab

# Lab exercises: UART to PWM

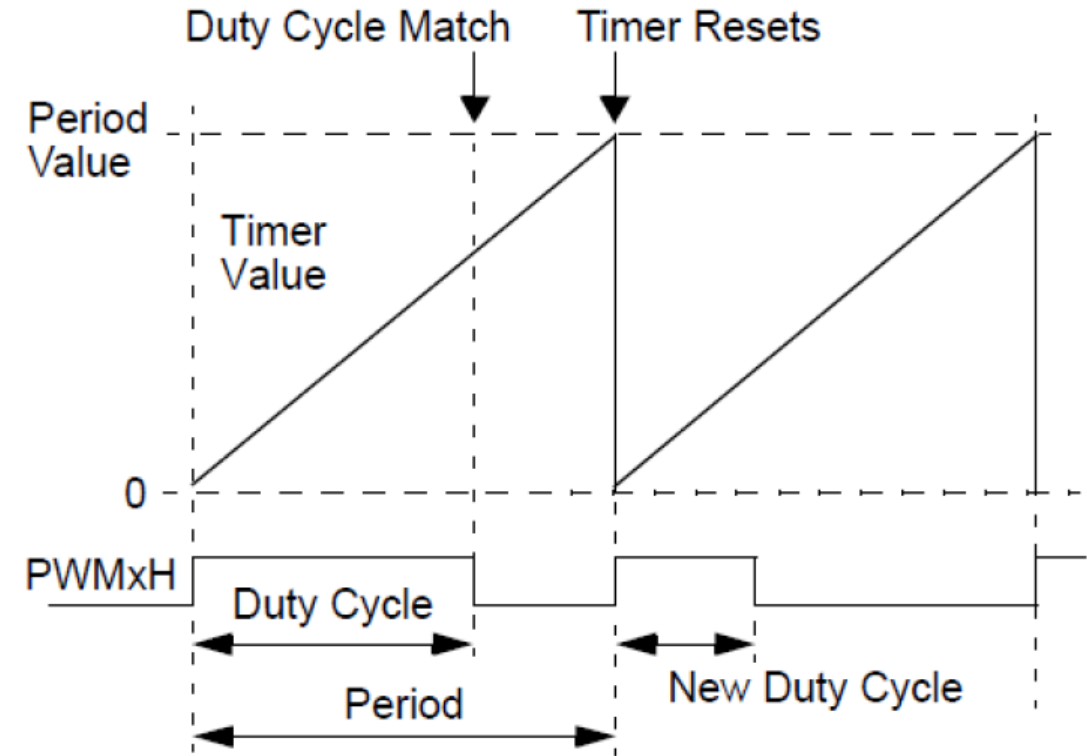
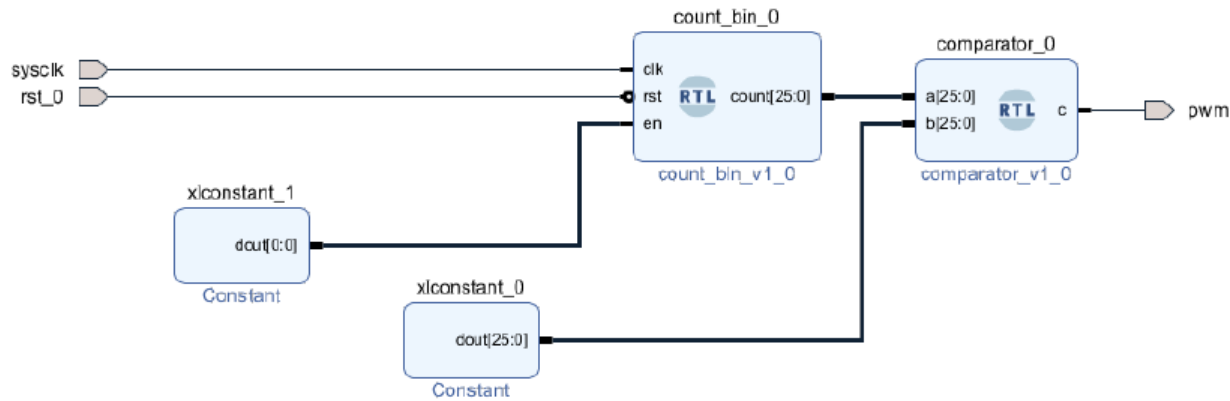
- The exercise aims to control the illumination of an LED using data coming from the UART, that represents the illumination value, and converts it into a PWM signal that will be sent to the LED.




- Source: <https://github.com/DIII-SDU-Group/EmbeddedSystems/tree/main/UART2PWM>



# Notes: PWM



# Notes: UART Overall system



- The UART solo code is in itsLearning

```

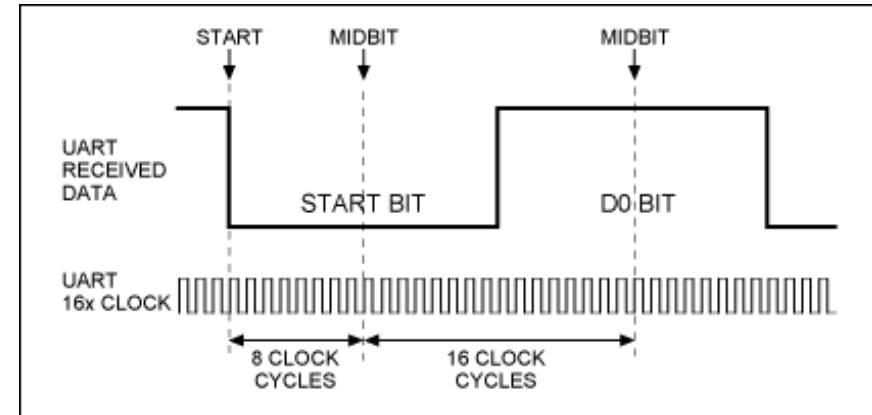
LIBRARY ieee;
USE ieee.std_logic_1164.all;
USE ieee.std_logic_arith.all;

ENTITY rx_mod IS
  PORT(
    clk      : IN      std_logic;
    rst      : IN      std_logic;
    sin      : IN      std_logic;
    data_out : OUT     std_logic_vector (7 DOWNTO 0);
    intr     : OUT     std_logic);
END rx_mod ;

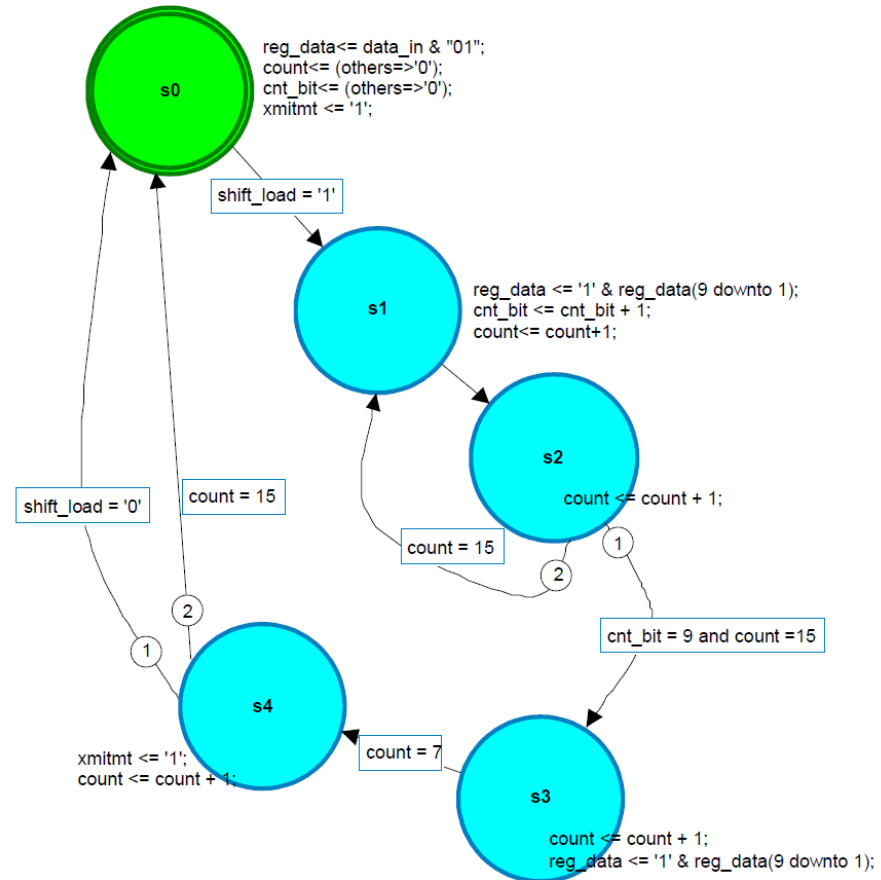
-- hds interface_end
ARCHITECTURE rtl OF rx_mod IS
  signal rxreg: std_logic_vector(9 downto 0);
  signal count: unsigned (3 downto 0);
  signal rxmt: std_logic;
  signal rxin,start_flag: std_logic;
begin
  process (clk, rst)
  begin
    if (rst = '0') then
      count <= (others => '0');
      rxmt <= '1';
      rxreg <= (others => '1');
      intr <= '0';
      rxin <= '1';
      start_flag<='0';
    elsif (rising_edge(clk)) then
      rxin<=sin;
      if (rxmt = '1' and rxin = '0') then
        count <= (others => '0');
        rxmt <= '0';
        rxreg <= (others => '1');
        start_flag<='0';
      elsif (count = 7 and rxmt = '0' and rxin = '0' and start_flag='0') then
        rxreg <= rxin & rxreg(9 downto 1);
        count <= (others => '0');
        start_flag<='1';
      elsif (count = 15 and rxmt = '0') then
        rxreg <= rxin & rxreg(9 downto 1);
        count <= count + 1;
      else
        count <= count + 1;
      end if;
      if (rxmt = '0' and rxreg(9) = '1' and rxreg(0) = '0') then
        intr <= '1';
        rxmt <= '1';
      else
        intr <= '0';
      end if;
    end if;
  end process;
  data_out <= rxreg(8 downto 1);
END rtl;

```

# Receiver



# Transmitter



- Default values are 0